



(a) Direct time-series-to-LLM alignment

$$p(y | X, q)$$




(b) Chart-based VLM inference

$$p(y | X, q)$$



(c) Numerical correction of visual tokens (ours)

$$p(y | X, q)$$



Frozen Trainable