



Soofi S pre-training dynamics


(a) Learning-rate schedule



(b) Language-modeling loss



(c) Gradient norm



(d) Training throughput

